

Profiling and Performance Improvement



Objectives

> After completing this module, you will be able to:

- >> Describe what profiling is and how it works
- >> Use the SDK profiling perspective
- >> Use profiling reports to evaluate software efficiency
- >> Discuss software tradeoffs to hardware
- >> Describe the function of the gprof tool
- >> List methods of improving performance

Outline

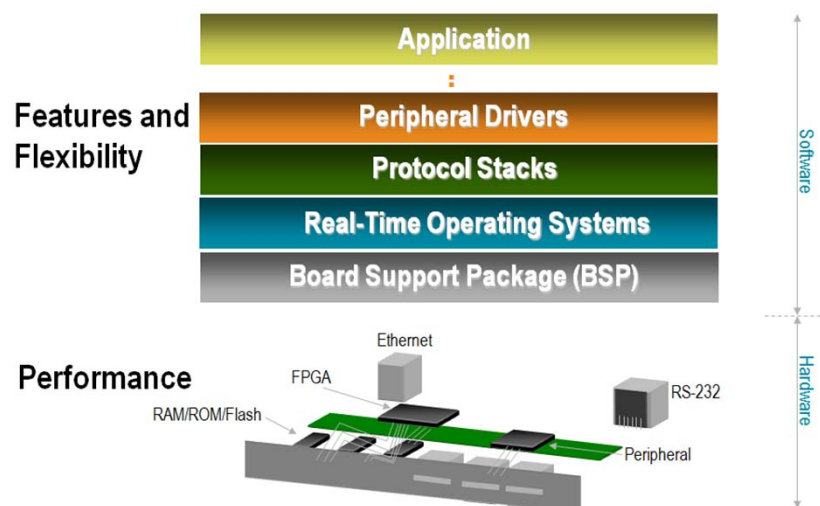
- > *Introduction*
- > Software Profiling
- > Profiling in SDK
- > Performance Improvement
- > Summary

Profiling and Performance 18-3

© Copyright 2018 Xilinx



Embedded Systems



Profiling and Performance 18-4

© Copyright 2018 Xilinx



Hardware and Software Partitioning

> Determine the software "critical path" by profiling

- >> Profiling measures where the CPU is spending its cycles on a function-by-function or task-by-task basis
- >> Similar to timing analysis in hardware
- >> Informs the system designer which software routine may be a candidate to hardware-accelerate

> Functions can be rewritten to improve efficiency in a number of ways

- >> Implementation in assembly code rather than C
- >> Writing faster C code, for example limit pointer use

What is Profiling?

> Profiling is an analysis of software performance

- >> Where routine time is being spent
- >> How many times functions are being called
- >> Included tool in SDK
- >> Which algorithms to consider moving to hardware

> Results in two useful formats

Name (location)	Time	Calls	Time/Call	%Time
Summary	403.599ms			100.0%
MAD_F_MUL_28	189.99ms	102400	1.846us	46.85%
MAD_F_MUL_28 (lab5.c:62)	96.299ms			23.86%
MAD_F_MUL_28 (lab5.c:61)	87.900ms			21.79%
MAD_F_MUL_28 (lab5.c:63)	4.900ms			1.21%
__muldd	107.799ms			26.71%
dct32	101.199ms	100	1.11ms	25.07%
dct32 (lab5.c:77)	45.699ms			11.32%
dct32 (lab5.c:76)	45.300ms			11.22%
dct32 (lab5.c:76)	9.100ms			2.25%
dct32 (lab5.c:71)	499.999us			0.12%
dct32 (lab5.c:74)	499.999us			0.12%

Samples per function: How much time is spent in each routine

Name (location)	Time	Calls	Time/Call	%Time
Summary	403.599ms			100.0%
MAD_F_MUL_28	189.99ms	102400	1.846us	46.85%
parents	101.199ms	102400	988ns	25.07%
dct32 (lab5.c:77)	101.199ms	102400	988ns	25.07%
__muldd	107.799ms			26.71%
dct32	101.199ms	100	1.11ms	25.07%
children	189.99ms	102400	1.846us	46.85%
MAD_F_MUL_28 (lab5.c:61)	189.99ms	102400	1.846us	46.85%
parents	2.300ms	100	23.0us	0.57%
main (lab5.c:117)	2.300ms	100	23.0us	0.57%
_exit	3.199ms			0.79%
main	2.300ms	0		0.57%
children	101.199ms	100	1.11ms	25.07%
dct32 (lab5.c:71)	101.199ms	100	1.11ms	25.07%

Function call graph: Which routine call, which function, and how many times

Software Profiling



© Copyright 2018 Xilinx

How Does Profiling Work?

- > **Hardware/software intrusive**
 - >> Requires a hardware timer
 - >> Requires a dedicated area in memory
 - >> Executable is modified with profiler routines
- > **A dedicated hardware timer interrupts the processor at a fixed interval**
 - >> The interrupt routine keeps track of the program counter at each interrupt
 - >> A histogram of PC locations is kept in profile RAM
 - >> Interrupt interval time is programmable
- > **Every function call in the software application is annotated by the compiler to track which functions are being called**

Coding Style Can Impact Profiling

- > **Effective profiling is based on how much time is spent in functions, and how often they are called**
 - >> If your code is just a fall-through main, profiling is not useful because 100 percent of execution time will be in *main()* with no calls to other functions
 - >> Carefully architect the application with a structured architecture by using functions
 - >> Compiler does not consider *macros* as functions – the macro will be expanded and treated as in-line code
 - >> Separate algorithms logically into functions that will help you analyze the flat profile view
 - >> Think ahead when architecting code—Is this algorithm a candidate for implementing in programmable logic?

Profiling Procedure

- > **Set profiling for the BSP**
 - >> Enable software intrusive profiling
 - >> Enable the -pg option
- > **Set profiling for the application**
 - >> Enable the compiler for profiling with the -pg option
 - >> Configure the profiler memory
 - >> Set the interrupt frequency and *bin* value
- > **Compile, link, and generate the ELF executable**
- > **Download the executable into a hardware or software simulator**
- > **Run the software application until completion or for an "amount of time"**
- > **Execute the GNU *gprof* tool to view the generated profile report**



Profiling in SDK



© Copyright 2018 Xilinx

Configuring the Software Platform Settings

- > Select Xilinx Tools > Board Support Package Settings
- > Select standalone
 - >> Enable software profiling
- > Select ps7_cortexa9_0
 - >> Add -pg to the Value column for the extra_compiler_flags option



Profile Configuration: Create a Run Configuration

- > If any of the embedded design resides in programmable logic, download the bitstream to the programmable logic
 - >> Select Xilinx Tools > Program FPGA
- > Select Run > Run Configurations and create a new configuration
 - >> Give appropriate name
 - >> Select the elf file that was compiled with `-pg`

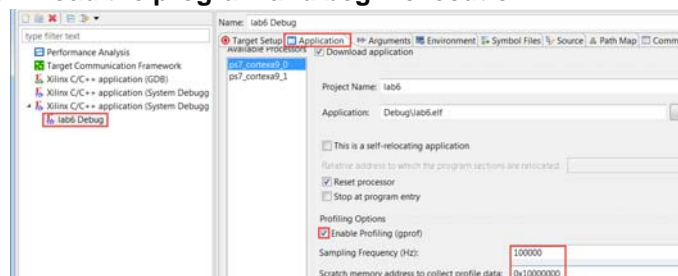
Profiling and Performance 18-13

© Copyright 2018 Xilinx



Set Profile Option in Run Configuration

- > In the Application tab
 - >> Enable profiling
 - >> Set the sampling frequency at which the timer will interrupt
 - Higher speed will require more memory but will give a finer resolution
 - >> Set the location of RAM that the profiler can use
 - make sure that the software application is not using this memory
- > Click Run to download the program and begin execution



Profiling and Performance 18-14

© Copyright 2018 Xilinx

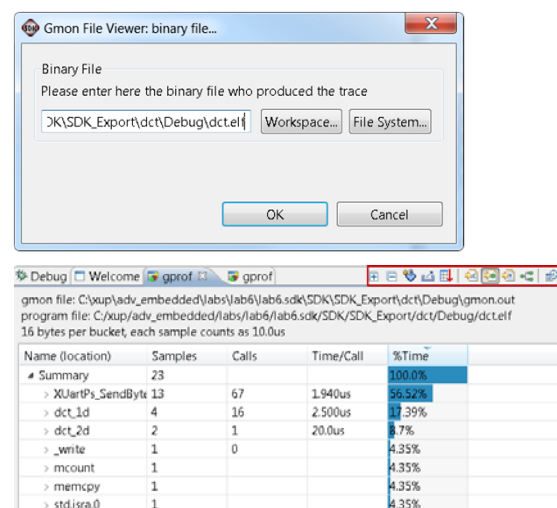


Profiling Reports

- > **Profile scratch memory is populated with statistics while the program is executing**
 - >> Intrusive profiling routines and the fixed interval timer interrupt use this memory
 - >> Stored in *gmon.out* upon completion or execution halt
- > **The *gprof* tool is launched by the tools after execution completion**
 - >> *gprof* reads *gmon.out* and assembles the information into a user configurable report

Viewing Profiling Reports: Launching *gprof*

- > **Double-click *gmon.out* to launch *gprof***
- > **Point to executable ELF; usually selected by default**
- > ***gprof* report launches**
- > **Report toolbar control report options and view capabilities**
 - >> Sort samples per file
 - >> Sort samples per function
 - >> Sort samples per line
 - >> Display function call graph
 - >> Switch sample/time



Profiled Output in SDK



Name (location)	Samples	Calls	Time/Call	%Time
Summary	23			100.0%
??	23			100.0%
> XJartPs_SendE 13	67		1.940us	56.52%
> dct_1d	4	16	2.500us	17.39%
> dct_2d	2	1	20.0us	8.7%
> _write	1	0		4.35%
> mcount	1			4.35%
> memcpy	1			4.35%
> std.isra.0	1			4.35%

1: Sort Samples per File

Name (location)	Samples	Calls	Time/Call	%Time
Summary	23			100.0%
> XJartPs_SendByte 13	67		1.940us	56.52%
> dct_1d	4	16	2.500us	17.39%
> dct_2d	2	1	20.0us	8.7%
> _write	1	0		4.35%
> mcount	1			4.35%
> memcpy	1			4.35%
> std.isra.0	1			4.35%

2: Sort Samples per Function

Name (location)	Samples	Calls	Time/Call	%Time
Summary	23			100.0%
> XJartPs_SendByte 13				56.52%
> dct_1d (??-1)	4			17.39%
> dct_2d (??-1)	2			8.7%
> _write (??-1)	1			4.35%
> mcount (??-1)	1			4.35%
> memcpy (??-1)	1			4.35%
> std.isra.0 (??-1)	1			4.35%
> Xil_In32 (??-1)	0			0.0%
> Xil_Out8 (??-1)	0			0.0%

3: Sort Samples per Line

Name (location)	Samples	Calls	Time/Call	%Time
Summary	23			100.0%
> XJartPs_SendByte 13	67		1.940us	56.52%
> dct_1d	4	16	2.500us	17.39%
> dct_2d	2	1	20.0us	8.7%
> _write	1	0		4.35%
> mcount	1			4.35%
> memcpy	1			4.35%
> std.isra.0	1			4.35%
> IRQInterrupt	0	0		0.0%

4: Display Function Call Graph

Profiling and Performance 18-17

© Copyright 2018 Xilinx

XILINX

Profiling Report Options



> Gprof report options allow report view flexibility and export

1

Show/Hide columns

☒ Name (location)
☒ Samples
☒ Calls
☒ Time/Call
☒ %Time

Select All

Deselect All

OK

Cancel

2

Export to CSV

Output file

CSV Configuration

CSV Separator: ;

☒ Expand All
☒ Export Hidden Columns

Tree sort: not drawing

☒ Report: sorted tree

Child Link: 1

No-child Link

Child Marker: --

Last child Marker: --

Leaf Marker:

Non-leaf Marker:

Profile:

Wineexecby:003_1:col_2

→ J21:00

→ J21:01

→ J21:02

→ J21:03

→ J21:04

→ J21:05

Restore Defaults

OK

Cancel

3

Sorting

Sort by:

1. %Time

2. Name (location)

3. Samples

4. Calls

Ascending

Descending

Restore Defaults

OK

Cancel

4

Switch time<->Samples

Name (location)	Samples	Calls
Summary	23	
> XJartPs_SendByte 13	67	
> dct_1d	4	16
> dct_2d	2	1
> _write	1	0

Name (location)	Time	Calls
Summary	230.0us	
> XJartPs_SendByte 130.0us	67	
> dct_1d	40.0us	16
> dct_2d	20.0us	1
> _write	10.0us	0

1. Show/hide columns

2. Export to CSV

3. Sorting

4. Switch time<->Samples

Profiling and Performance 18-18

© Copyright 2018 Xilinx

XILINX

9

Performance Improvement



© Copyright 2018 Xilinx

Task Implementation Decision

> Keep it in software

- >> Not in critical path
- >> Enough "free" cycles
- >> Easier to code in software than in hardware
 - Uses math library functions
- >> NEON co-processor
 - Supports integer vector operations
 - Single floating-point operations

> Move to hardware

- >> Programmable logic co-processor
 - Customized to user's needs
 - Excellent for iterative and pipelined processing
- >> Add soft core processor in PL
 - Both Cortex-A9 and MicroBlaze processors can co-exist in the SoC

Software to Programmable Logic

- > **Slow software tasks can be accelerated by taking them to hardware**
 - >> Start with functions where the software spends most of its time
 - >> Consider the hardware implementation and if there is potential benefit implementing in hardware
- > **Many mechanisms**
 - >> Dual-port block RAM
 - >> Custom AXI peripheral
 - >> Code optimization: use of macros, increasing compiler optimization
 - >> Enabling caching if (by default) it is turned OFF

Using Block RAM

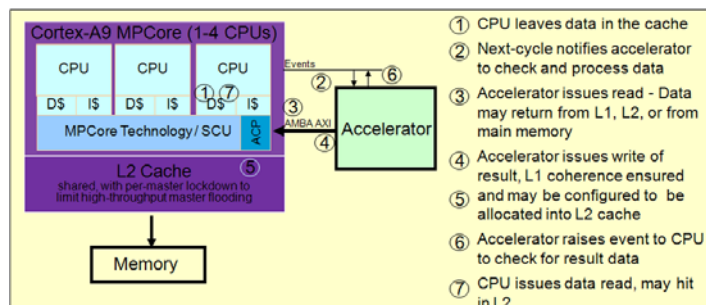
- > **Leverage the dual-port nature of Xilinx block RAM**
- > **Useful for data in block or frame format**
 - >> Video
 - >> 2D matrix maps
- > **Advantages**
 - >> Low silicon overhead
 - >> Fast and deterministic latency



Enhanced Accelerator SoC Integration

> ARM MPCore: accelerator coherence port (ACP)

- >> Sharing benefits of the ARM MPCore optimized coherency design
- >> Accelerators gain access to CPU cache hierarchy
- >> Compatible with standard un-cached peripherals and accelerators



Profiling and Performance 18-23

© Copyright 2018 Xilinx

XILINX

Summary

XILINX

© Copyright 2018 Xilinx

Summary

- > Profiling allows you to analyze the software and determine where the CPU's time is spent
- > Profiling can help you rearrange or rewrite the code or even help you consider if a function can be targeted to hardware
- > Enabling cache can improve performance
- > The gprof tool is used to generate a profiling report from collected statistics
- > A hardware timer and memory are required to use the profiling tool
- > Profiling in SDK is provided by the Standalone BSP as a GNU service

Adaptable.
Intelligent.